

AI Transformation Specialist – Technical Assignment

Duration: 2–3 Days

Role: AI Transformation Specialist (2–4 Years Experience)

Company: [Techculture.ai](#)

Overview

This assignment evaluates your practical ability to design and implement an end-to-end AI solution using Retrieval-Augmented Generation (RAG) and agentic AI concepts. You will build a small, functional AI assistant that demonstrates your understanding of GenAI, vector databases, LLM orchestration, and business-focused solution design.

Problem Statement

You are tasked with creating a **Service Information Assistant** for a technology services company. This assistant will:

- Ingest and index company service documents (AI consulting, digital marketing, web/app development, FAQs, case studies)
- Answer user questions accurately using a RAG pipeline
- Demonstrate basic "agentic" behaviour by calling a simple tool or API when appropriate
- Provide source citations for transparency and trust

Scenario: A potential client or internal team member asks questions like:

- "What AI services do you offer?"
- "Do you have experience with fintech projects?"
- "What's the estimated cost for a 3-month digital marketing campaign?"

Your assistant should retrieve relevant information from documents and generate accurate, context-aware answers.

Core Requirements

1. Data Preparation & Document Ingestion

- Create or use 5–10 text sources (Markdown, TXT, or PDF) representing:
 - Service descriptions (AI consulting, digital marketing, development)
 - FAQ section
 - 1–2 case study examples or client scenarios
- If we don't provide real documents, create dummy content that simulates realistic service information
- Design the system to be scalable for real production documents

2. RAG Pipeline Implementation

Build a complete Retrieval-Augmented Generation pipeline:

- **Document Processing:** Split documents into semantic chunks (appropriate size for retrieval)
- **Embeddings:** Generate vector embeddings using a suitable model (e.g., OpenAI embeddings, sentence-transformers)
- **Vector Storage:** Store embeddings in a vector database (FAISS, Chroma, Pinecone, or similar)
- **Retrieval:** Implement semantic search to retrieve top-k relevant chunks for user queries
- **Generation:** Use an LLM (OpenAI API, Hugging Face, or other) to generate answers grounded in retrieved context

Key Feature: All answers must include **source citations** showing which document/section was used (e.g., "Source: services_overview.txt, Section: AI Consulting").

3. Agentic Behaviour (Tool Integration)

Implement at least **one tool** that your AI assistant can call autonomously:

Option A – Sentiment Analyzer:

- Classify user query sentiment (Positive/Negative/Neutral)
- Include sentiment context in the response when relevant

Option B – Pricing Estimator:

- Simple Python function that takes service type and duration
- Returns rough cost estimate range
- Agent decides when to call this based on query intent

Option C – Custom Tool:

- Design your own lightweight tool relevant to the business context
- Demonstrate autonomous decision-making about when to use it

You may use LangChain agents/tools framework or implement your own orchestration logic.

4. User Interface

Provide at least one of the following:

Option A – REST API (FastAPI):

- Endpoint: POST /ask
- Request body: {"question": "your question here"}
- Response: Answer + sources + tool outputs (if applicable)
- Include basic error handling

Option B – Interactive Interface:

- Simple Streamlit app or CLI interface
- User enters question and sees answer with sources
- Display tool usage when triggered

Technical Expectations

Required Stack

- **Language:** Python 3.8+
- **Core Libraries (suggested):**
 - LangChain or similar for RAG/agent orchestration
 - FAISS / Chroma / Pinecone for vector storage
 - OpenAI API / Hugging Face / other LLM
 - FastAPI or Streamlit (for interface)
- **Additional Tools:** Pandas, NumPy, NLTK/spaCy (as needed)

Code Quality Standards

- Clear project structure with separated modules
 - Configuration management (API keys, model settings in config file or environment variables)
 - Basic error handling and input validation
 - Commented code where logic is complex
 - Requirements.txt with all dependencies
-

Deliverables

Submit the following within 2–3 days:

1. Code Repository

- Complete project folder or GitHub repository link
- Suggested structure:
 - ingestion.py – Document loading and chunking
 - embeddings.py – Vector embedding generation
 - rag_pipeline.py – Retrieval and generation logic
 - agent.py – Tool integration and orchestration
 - api.py or app.py – Interface layer
 - config.py – Configuration settings
 - requirements.txt – Dependencies
- Include all dummy documents or clear instructions on generating them

2. README Documentation

Comprehensive [README.md](#) including:

Setup Instructions:

- Step-by-step guide to install dependencies
- How to configure API keys and environment
- Commands to run the application

Architecture Overview:

- High-level data flow diagram or bullet points:
 - Documents → Chunking → Embeddings → Vector Store
 - Query → Retrieval → LLM Generation → Response
 - Agent/Tool decision flow
- Technology choices and justification

Assumptions & Limitations:

- What assumptions did you make about the data or use case?
- Known limitations of your current implementation

Future Improvements:

- What would you add given more time?
- How would you scale this for production?

3. Sample Queries & Results Report

Create a separate document (Markdown or PDF) with:

- **8–10 example questions** covering different scenarios:
 - Service information queries
 - FAQ-style questions
 - Questions requiring tool usage
 - Edge cases or challenging queries
- For each query, include:
 - The question asked
 - System's answer
 - Sources cited
 - Tool outputs (if triggered)

- Your analysis: "Why this answer is good/acceptable" or "What could be improved"

Example format:

Query 1: "What AI consulting services do you provide?"

Answer: [System response with context]

Sources: services_overview.txt (AI Consulting section)

Analysis: Answer accurately reflects document content with proper citations. Could be enhanced with specific case study references.

Evaluation Criteria

We will assess your submission based on:

Criteria	What We Look For
RAG Implementation	Proper chunking, embedding, retrieval, and generation pipeline
Agentic Behaviour	Thoughtful tool integration with autonomous decision-making
Code Quality	Clean structure, modularity, error handling, documentation
Solution Design	Understanding of business context and practical applicability
Documentation	Clear README, architecture explanation, honest assessment
Testing & Examples	Comprehensive sample queries demonstrating system capabilities
Innovation	Creative approaches, additional features, or optimizations

Submission Guidelines

1. Complete the assignment within 2–3 days from receipt
 2. Submit via:
 - GitHub repository link (preferred – ensure it's public or add our reviewer)
 - OR ZIP file containing complete project
 3. Email submission to: [Insert HR contact email]
 4. Subject line: "AI Transformation Assignment – [Your Name]"
 5. Include:
 - Link or attachment to code
 - [README.md](#)
 - Sample queries report
 - Any additional notes or questions
-

Important Notes

- You may use any open-source libraries or frameworks
 - If using paid APIs (OpenAI), keep costs minimal – we understand token limits
 - Focus on demonstrating **understanding and practical skills** over perfection
 - If you face blockers, document your approach and reasoning in README
 - Authenticity matters – we value your own thinking over copied solutions
-

Questions?

If you have clarifying questions about the assignment, please reach out to:

Contact: [HR Name/Email]

Response Time: We'll respond within 24 hours

Good Luck!

We're excited to see your approach to building AI solutions. This assignment mirrors real work you'd do as part of our team – bridging technical implementation with business value.

Focus on demonstrating:

- Your technical depth in AI/ML and GenAI
- Your ability to design practical, end-to-end solutions
- Your communication through clean code and clear documentation

We look forward to reviewing your submission!

Techculture.ai – AI Transformation Team